

# Introduction to Linux containers

Application in scientific practice  
Docker and Singularity

Toni Hermoso Pulido  
Bioinformatics Unit  
CRG, Barcelona

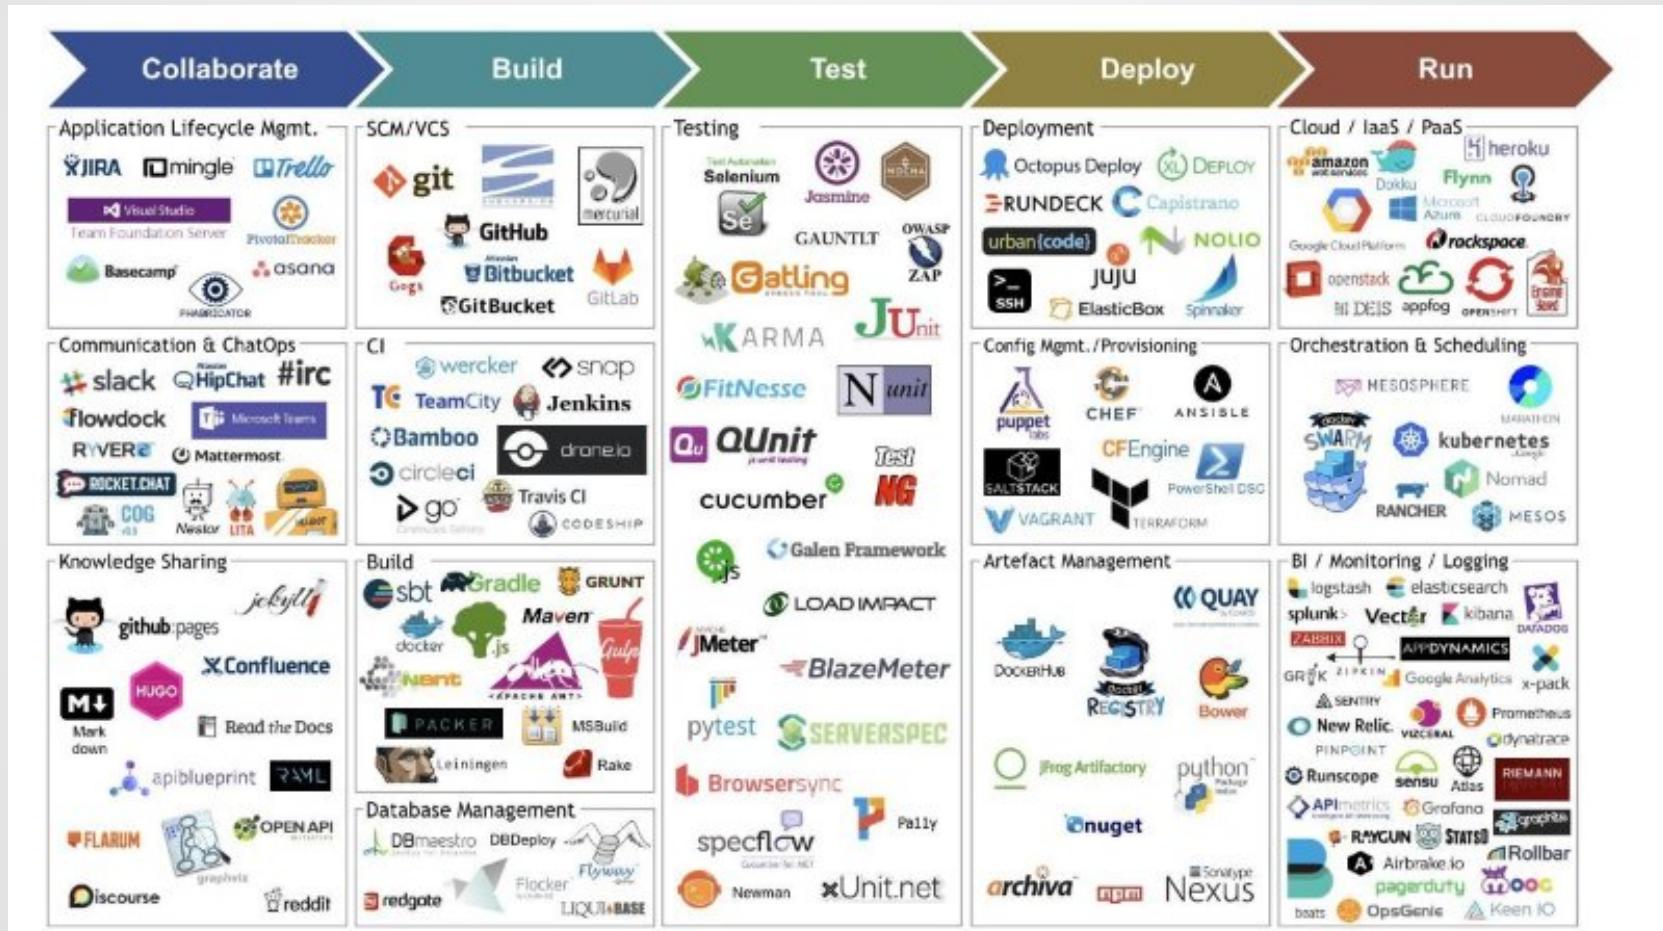
**Recipes and instructions:**

<https://github.com/biocorecrg/containers-course>

# DevOps

- Software Engineering Culture
- Software Development + Software Operations
- Automate and monitor

# DevOps



<https://hostadvice.com/blog/devops-toolbox-jenkins-ansible-chef-puppet-vagrant-saltstack/>

# Containers



Containers in New Jersey

# Virtualisation

- Abstraction of physical hardware
- Depends on hypervisor (software)
- Do not confuse with hardware emulator
- Enable virtual machines
  - Every virtual machine with an Operating System

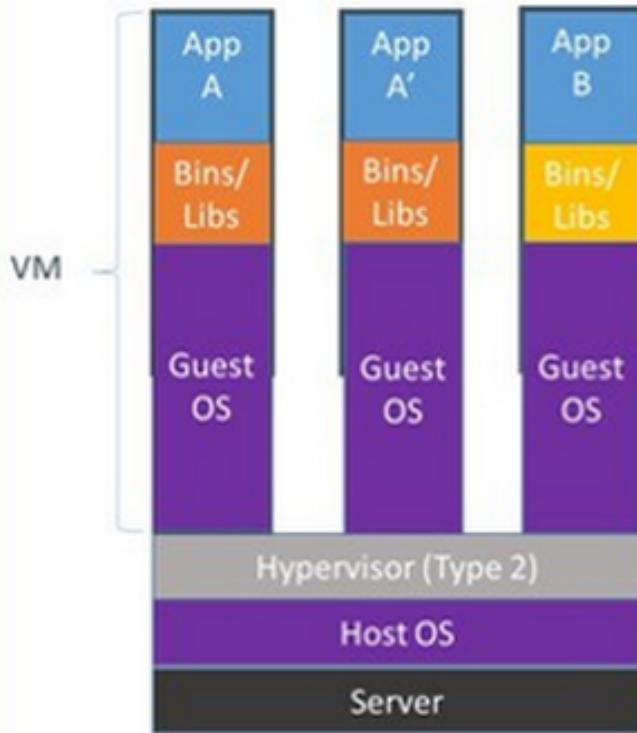
# Containerisation

*aka Lightweight virtualisation*

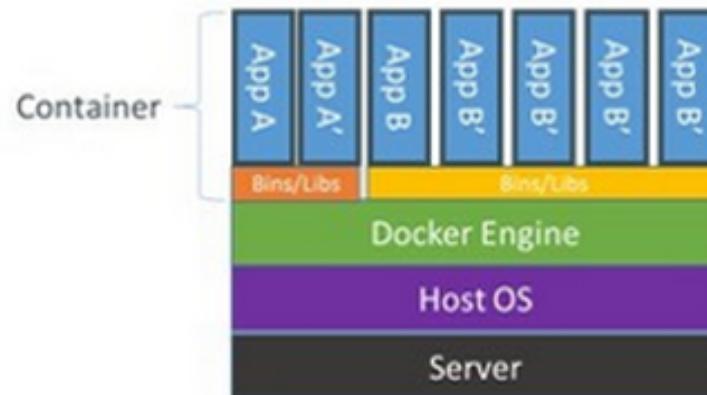
- Abstraction of application layer
- Depends on host kernel (Operating System)
- Application and dependencies bundled all together

# Virtual machines vs containers

## Containers vs. VMs



Containers are isolated, but share OS and, where appropriate, bins/libraries



# Virtualisation

## Pros and Cons

- **PRO:** Very similar to a full OS
- **PRO:** With current solutions, high OS diversity
- **CON:** Need of more space and resources
- **CON:** Slower than containers
- **CON:** Not as good automating

# Containerisation

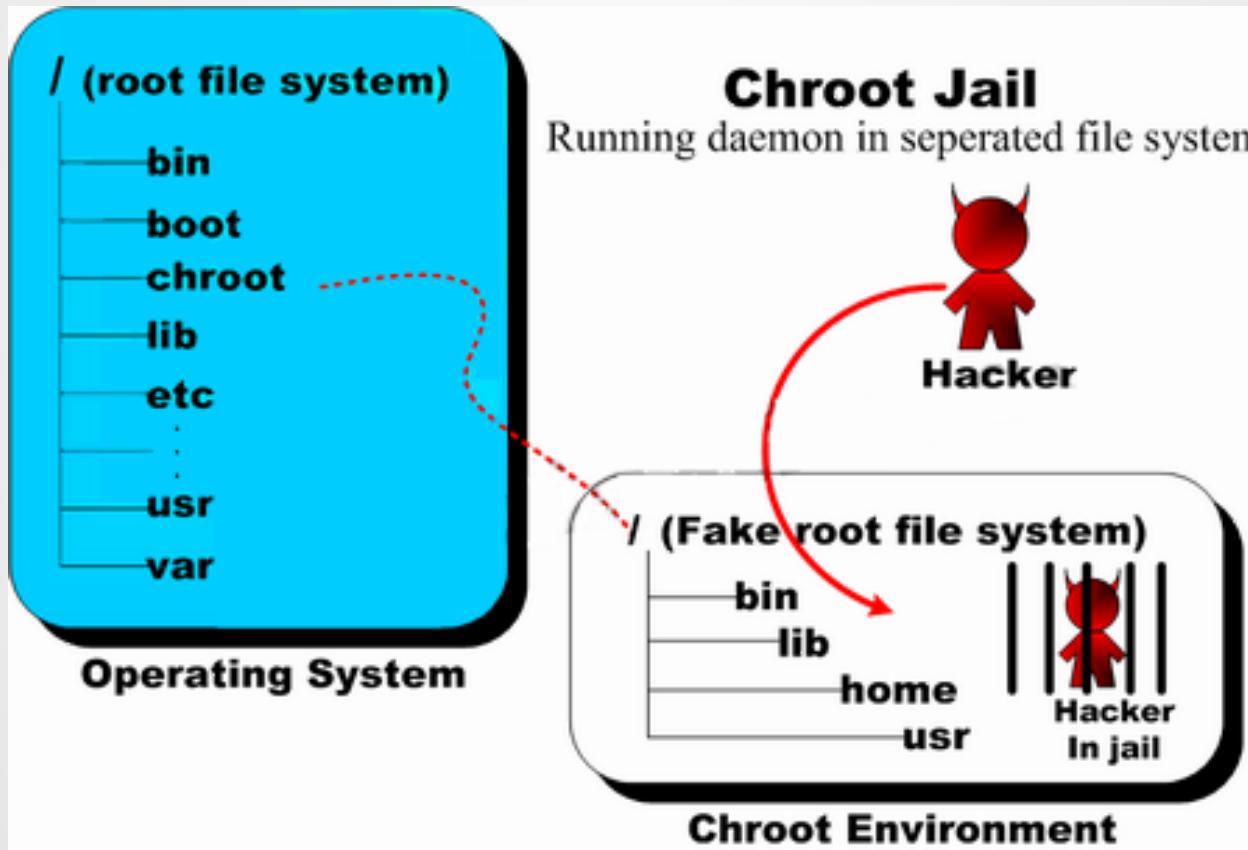
## Pros and Cons

- **PRO:** Faster
- **PRO:** No need of full OS installation. Less space.
- **PRO:** Current solutions allow easier distribution of recipes. More portability
- **PRO:** Easier automation
- **CON:** Some cases might not be exactly the same as a full OS
- **CON:** With current solutions, still less OS diversity

# History of containers

# chroot

- chroot jail (BSD jail) - First concept 1979
- Notable use in SSH and FTP servers
- Honeypot, recovery of systems, etc.



# Additions in Linux kernel

- [cgroups](#) (control groups), before 'process containers'
  - isolates resource usage (CPU, memory, disk I/O, network, etc.) of a collection of processes
- [Linux namespaces](#)
  - one set of kernel resources restrict to one set of processes

# Additions in Linux kernel



Linux Containers



liblxc



Docker 1.10 and later



runC

runC

runC

containerd-shim

containerd-shim

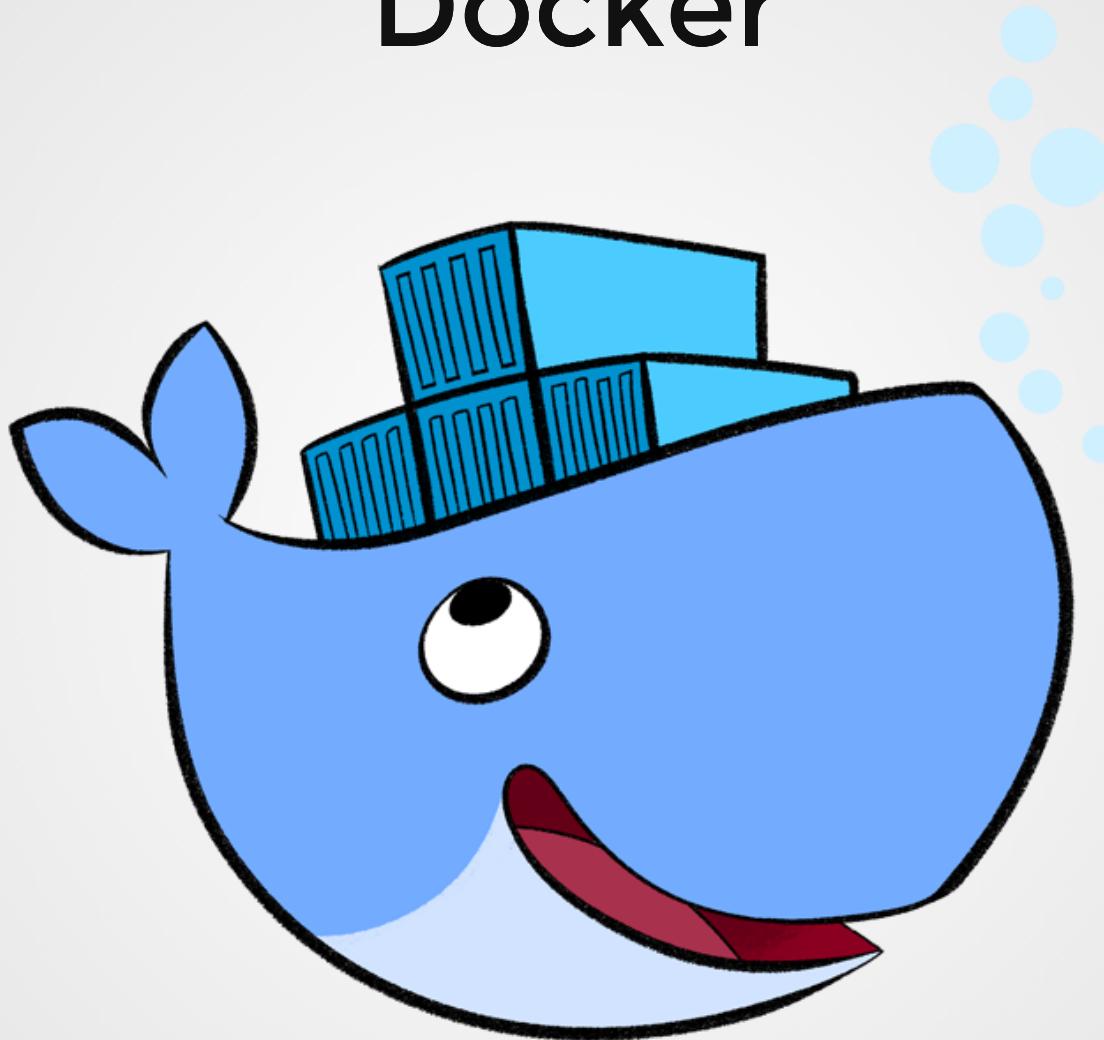
containerd-shim

containerd

Docker Engine



# Docker



# Docker

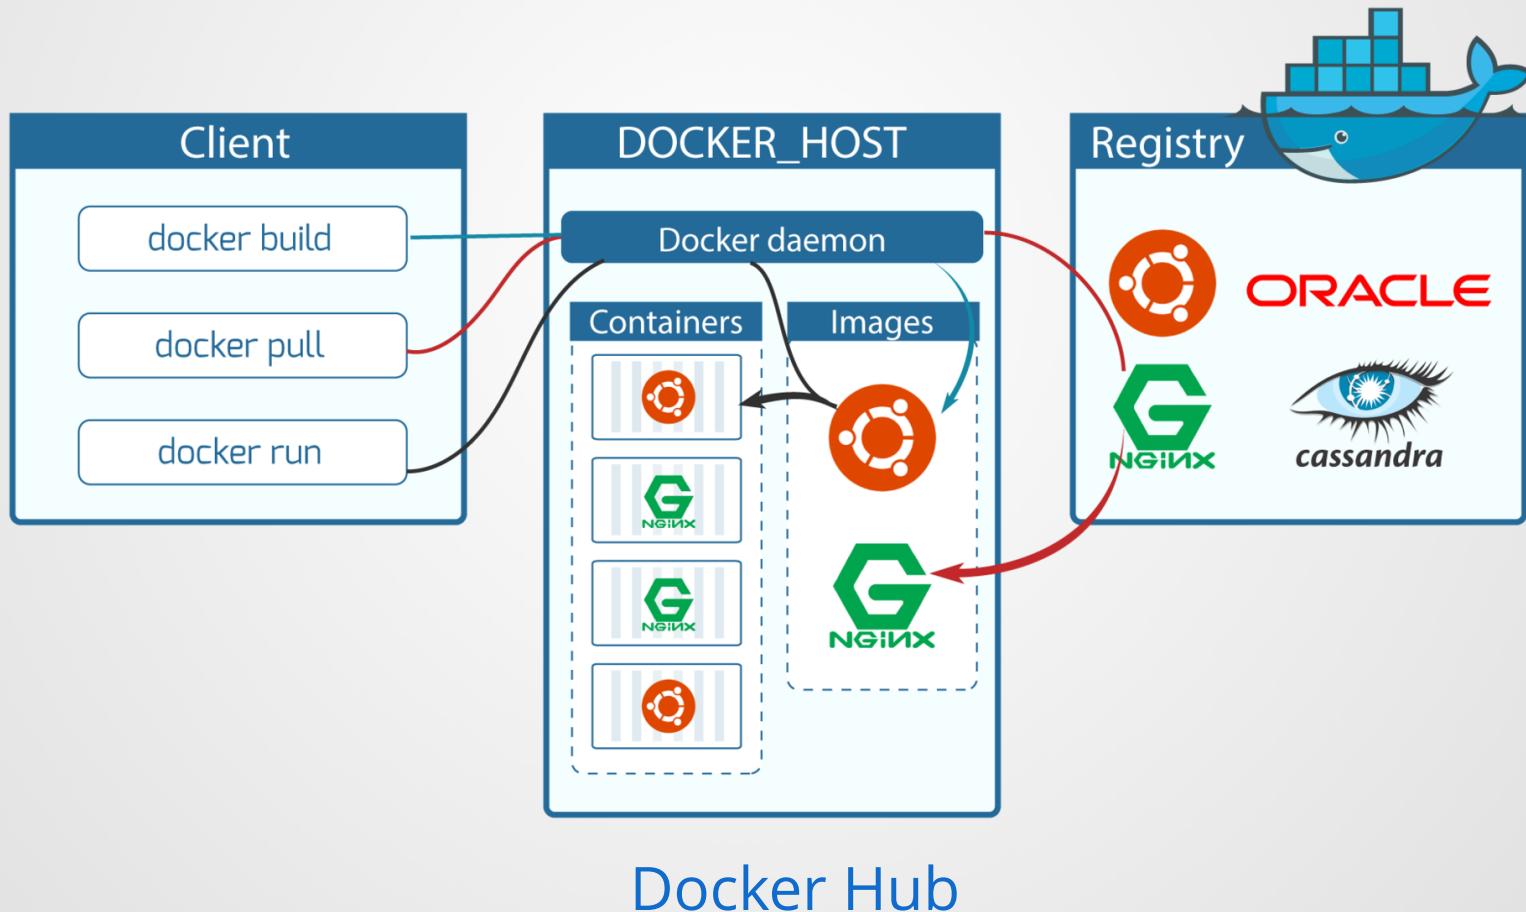
- Platform for developing, shipping, and running applications
- Infrastructure as application/code
- Established [Open Container Initiative](#)

As a software:

- [Docker Community Edition](#)
- Docker Enterprise Edition

# Docker architecture

## DOCKER COMPONENTS



# Docker image

- Read-only templates.
- Containers are run from them
- Images are not run
- Images have several layers

# Docker registry and Docker hub

- Images are stored locally
- They can also be shared in a registry
- Main Public one: [Docker hub](#)

Examples:

<https://hub.docker.com/u/biocorecrg/>

<https://github.com/CRG-CNAG/docker-debian-perlbrew>

# Docker registry commands

download image

```
$ docker pull debian  
$ docker pull debian:stretch  
$ docker pull centos:7  
$ docker pull rocker/r-ver:3.5.1  
$ docker pull continuumio/miniconda:4.5.4
```

Debian

CentOS

Rocker R

Miniconda

# Docker container

- Generated from an image (template)
- Image: read-only
- Container: read-write
- Can be converted into image
  - docker commit
  - docker import & docker export
- 1 image -> n diverse containers
  - Diversity:
    - Volumes / Mounting points
    - Different data or configs
    - Different exposed ports

# Run container

```
$ docker run debian /bin/echo 'Hello world!'
$ docker run debian:jessie /bin/echo 'Hello world!'
$ docker run debian:jessie cat /etc/issue
$ docker run debian:stretch cat /etc/issue
```

# Run container as daemon

Run daemon

```
$ docker run -d debian:stretch tail -f /dev/null  
$ docker run -d --name mycontainer debian:stretch tail -f /dev/null
```

List running containers

```
$ docker ps
```

Log and info of container

```
$ docker log mycontainer  
$ docker inspect mycontainer
```

Actions on containers

```
$ docker stop mycontainer  
$ docker start mycontainer  
$ docker restart mycontainer
```

# More running on container

Execute on a running container

```
$ docker exec mycontainer /bin/echo 'Bye, moon!'
```

Run a container interactively (from the beginning). Stops when exiting

```
$ docker run -ti --name mycontainer2 debian:stretch /bin/bash
```

Execute on a running container interactively

```
$ docker exec -ti mycontainer /bin/bash
```

# Other Docker commands

remove container, remove image

```
$ docker rm mycontainer  
$ docker rmi myimage
```

list all containers (even stopped ones)

```
$ docker ps -a
```

save image as tar

```
$ docker save myimg > myimg.tar
```

```
$ docker export mycontainer > mycontainer.tar
```

```
$ docker import mycontainer.tar mycontainer:latest
```

import tar as image

# Docker image - Building

- Can be built from existing base images
  - E.g. : Ubuntu, Apache
- Any modification from a base image is a new layer
- Base images can be generated with tools such as [Debootstrap](#)
  - Reminds of creation of a bootable (USB) Image

# Docker image - Instructions

- Recipe file:
  - Dockerfile
- Instructions
  - *Every instruction generates an image layer*
  - FROM: use a base image (notice tag)
  - ADD, COPY: add files to image filesystem
  - RUN: execute command in image
  - ENV, ARG: Run and build environment variables
  - CMD, ENTRYPOINT: Command to execute when generated container starts

# Docker image - Instructions

**ADD vs COPY**

in Stack Overflow

# Docker image - Instructions

- Instructions
  - USER: User to run a following process (fallback: *root*)
  - WORKDIR: Location where following processes will be executed
  - VOLUME: Definition of mounting point (with host)
  - EXPOSE: Ports to be accessible from the container

Reference  
Best Practices

# Docker service mappings

## Ports and volumes (directories/files)

**host:container**

```
docker run -d -v /db/test:/blastdb -p 8089:8081 --name myblast blastww
```

# Docker image

## Services approaches

### ENTRYPOINT vs CMD

- Entrypoint script with or without appended command
- tini / --init
- supervisord

# Containers in biosciences

Biocontainers: repository of bioinformatics containers  
<http://biocontainers.pro>

<https://dx.doi.org/10.1093/bioinformatics/btx192>

nf-core  
<https://hub.docker.com/r/nfcore/>

# Docker registry commands

download image

```
$ docker pull biocorecrg/debian-perlbrew:stretch
```

upload image

```
$ docker push myusername/mysexyimage
```

# Docker automated builds

Reference

# Docker Compose

Reference

`docker-compose.yml`

# Other Docker tools

**Portainer**

**Dive**

# Singularity

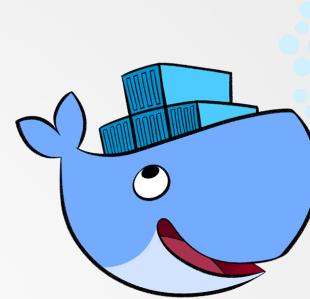
containers for HPC



<https://www.sylabs.io/singularity/>

<https://doi.org/10.1371/JOURNAL.PONE.0177459>

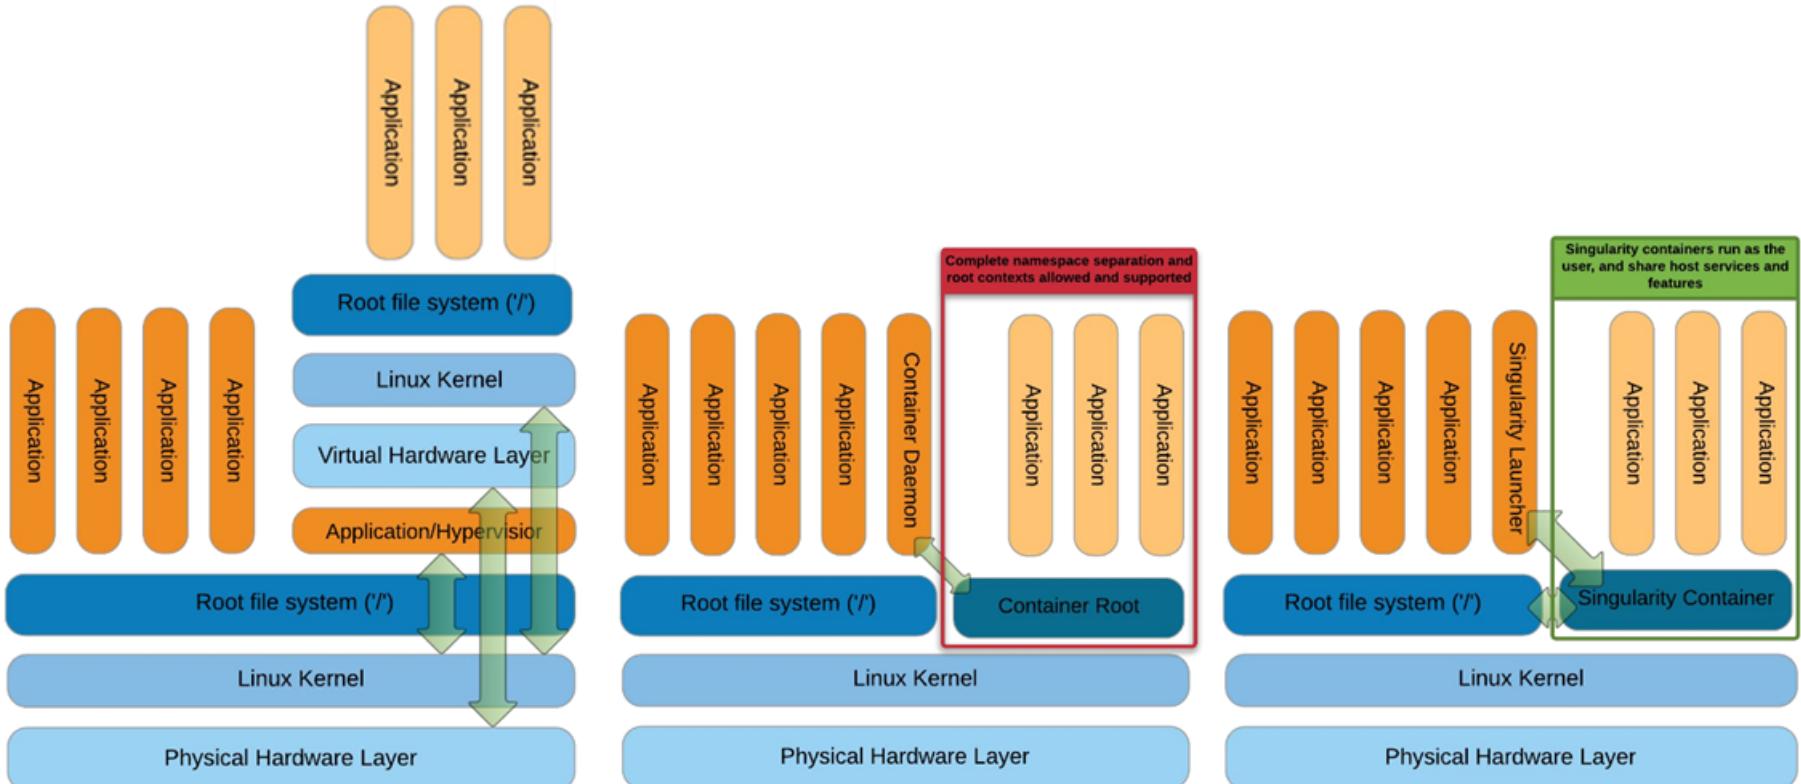
# Singularity vs Docker



Summarising

- Docker -> Microservices
- Singularity -> HPC

# Singularity architecture



General VM  
eg ESXi

General Container  
eg Docker

HPC Container  
Singularity

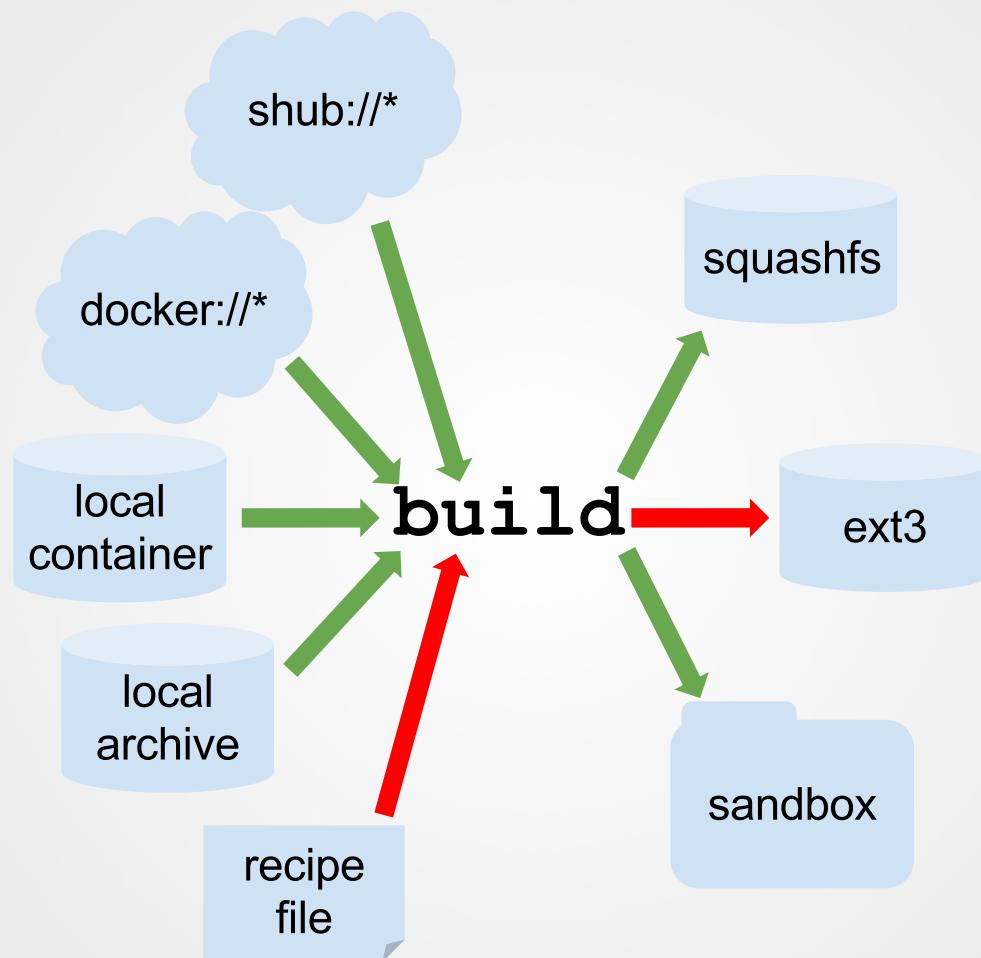
# Singularity - Strengths

- No dependency of a daemon
- Can be run as a simple user
- Image/container is a file (or directory)
  - More easily portable
- Two type of images
  - Read-only (production)
  - Writable (development)

# Singularity - Weaknesses

- At the time of writing only good support in Linux
  - *Not a big deal in HPC environments, though*
- For some uses you might need root account (or sudo)

# Singularity - build



As of 2.4.x version

<http://singularity.lbl.gov/docs-build-container>

# Singularity - build

Build read-only image from Docker

```
$ singularity build debian.sif docker://debian:stretch
```

Build writable directory from Docker

```
$ sudo singularity build --sandbox debiandir docker://debian:stretch
```

# Singularity - run

Execute a command

```
$ singularity exec debian.sif /bin/echo 'Hello world'  
$ singularity exec debian.sif env
```

Execute a command (with clean environment)

```
$ singularity exec -e debian.sif /bin/echo 'Hello world';  
$ singularity exec -e debian.sif env
```

Execute a shell

```
$ singularity shell debian.sif
```

Execute defined runscript (parameters can be used)

```
$ singularity run debian.sif
```

# Private Docker images

[https://www.sylabs.io/guides/3.2/user-guide/singularity\\_and\\_docker.html](https://www.sylabs.io/guides/3.2/user-guide/singularity_and_docker.html)

```
sudo singularity build privateimg.sif docker-daemon://privateimg:latest
```

```
sudo singularity build privateimg.sif docker-archive://privateimg.tar
```

# Singularity recipes

```
BootStrap: debootstrap
OSVersion: stretch
MirrorURL: http://ftp.fr.debian.org/debian/
Include: curl

%environment
    BLAST_PROGRAM=blastp
    BLASTDB=/blastdb
    export BLAST_PROGRAM BLASTDB

%post
    BLAST_VERSION=2.7.1

    cd /usr/local; curl --fail --silent --show-error --location --remote-name ftp://ftp.ncbi.nlm.nih.gov/blas
    cd /usr/local; tar zxf ncbi-blast-${BLAST_VERSION}+-x64-linux.tar.gz; rm ncbi-blast-${BLAST_VERSION}+-x64
    cd /usr/local/bin; ln -s /usr/local/ncbi-blast-${BLAST_VERSION}+/bin/* .

%labels
    Maintainer Biocorecrg
    Version 0.1.0

%runscript
    echo "Blast application!"
    exec $BLAST_PROGRAM "$@"
```

# Singularity recipes

```
BootStrap: docker
From: continuumio/miniconda:4.5.4

%environment
    PATH=/opt/conda/envs/blast-conda/bin:$PATH
    export PATH

%post
    PATH=/opt/conda/bin:$PATH
    export PATH
    conda env create -f /environment.yml && conda clean -a

%files
    environment.yml

%labels
    Maintainer Biocorecrg
    Version 0.1.0
```

# Singularity - volumes

```
singularity run -B /db/test:/blastdb blastwww.sif
```

# Singularity - services/instances

```
sudo singularity build blastmysql.sif Singularity
singularity run -B /db/test:/blastdb -B $(pwd)/config.json:/config/mysql.json blastmysql.sif
sudo singularity build mysql.sif Singularity.mysql
singularity exec -B /tmp/db:/var/lib/mysql mysql.sif mysql_install_db
singularity instance.start -B /tmp/db:/var/lib/mysql -B /tmp/socket:/run/mysqld mysql.sif mysql
singularity instance.list
singularity exec instance://mysql mysql -uroot -h127.0.0.1 -e "CREATE DATABASE blast; GRANT ALL PRIVILEGES
singularity instance.start -B /db/test:/blastdb -B $(pwd)/config.local.json:/config/mysql.json blastmysql.
singularity instance.list
singularity instance stop.blast
singularity instance.stop mysql
```

# Further reading

- The impact of Docker containers on the performance of genomic pipelines.
- Performance Evaluation of Container-based Virtualization for High Performance Computing Environments

# Sum-up: containers in science

- Maintainability
- Portability
- Reproducibility

# Play!

<https://github.com/biocorecrg/containers-course>